

CSCI 4140 – Tutorial 7

Learning the basics of Node.js

Matt YIU, Man Tung ([mtyiucse](mailto:mtyiucse@gmail.com))

SHB 118

Office Hour: Tuesday, 3-5 pm

2015.03.05

Outline

- What is Node.js?
- Learning the basics of Node.js: Non-blocking I/O, HTTP
 - Exercises adapted from **learnyounode**:
<https://github.com/rvagg/learnyounode>

License of learnyounode

learnyounode is Copyright (c) 2013-2015 learnyounode contributors (listed above) and licenced under the MIT licence. All rights not explicitly granted in the MIT license are reserved. See the included LICENSE.md file for more details.


learnyounode builds on the excellent work by @substack and @maxogden who created stream-adventure which serves as the original foundation for **learnyounode**.

What is Node.js?

- An **open-source, cross-platform** runtime environment for **server-side** and **networking applications**
- Applications are written in **JavaScript**
 - Node.js uses **Google V8 JavaScript engine** to execute code
- Provide an **event-driven architecture** and a **non-blocking I/O API**
 - One process for all **concurrent connections**
 - Optimizes an application's **throughput** and **scalability**
 - For your information, Apache uses **process-/thread-based architecture**, which is relatively inefficient
 - A new process / thread is created per connection


What is Node.js: Event-driven architecture

Event Emitters



Event Queue

Event Loop
(single-threaded)



For those who have taken
CSCI/CENG 3150...



States



Event Handler

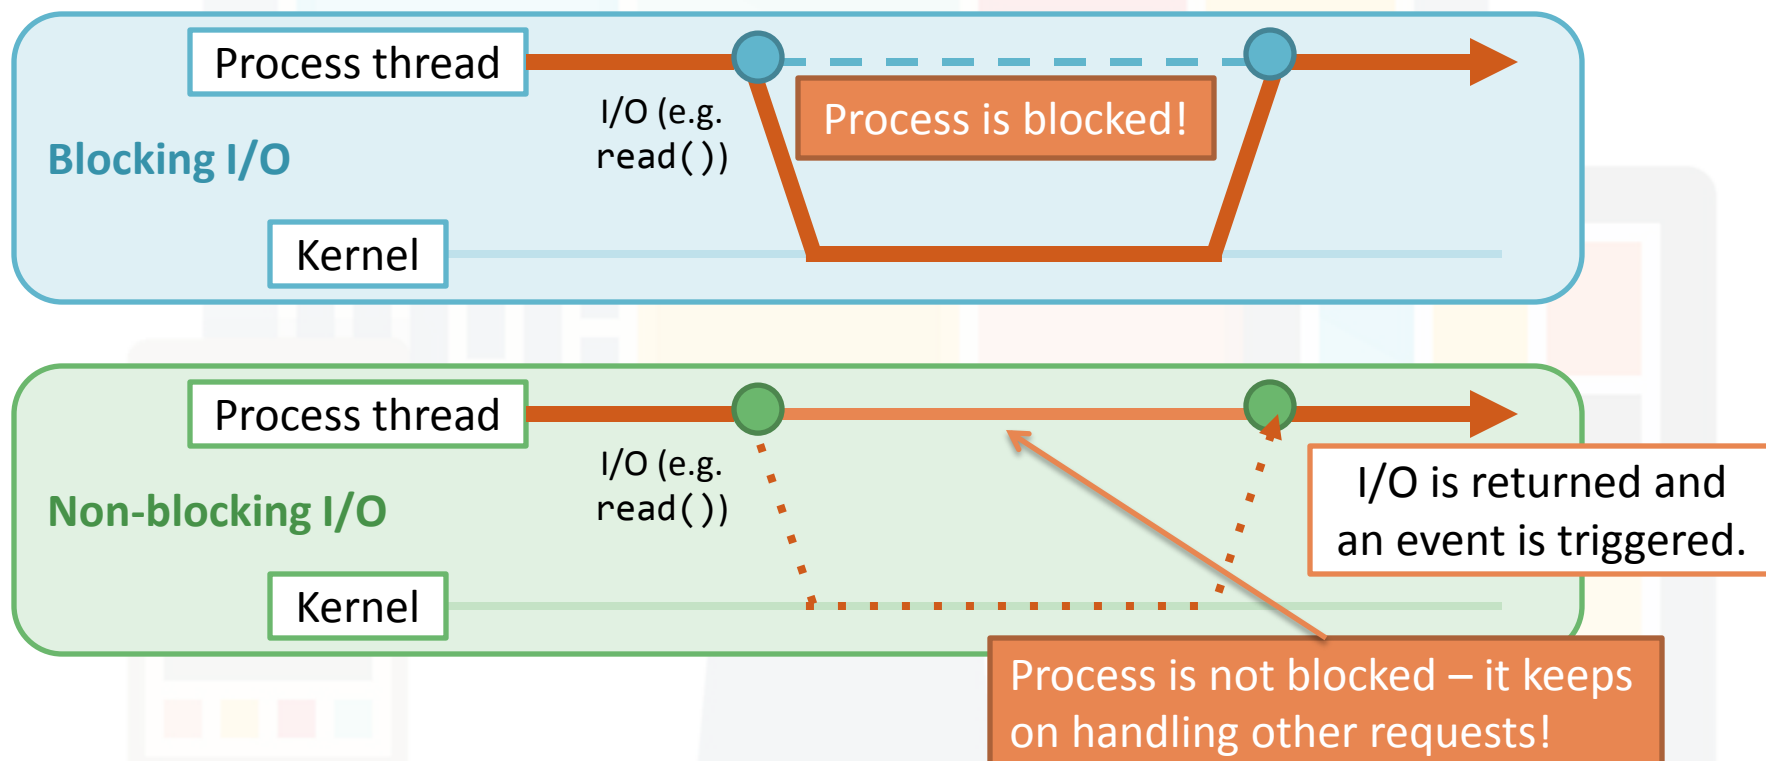
One thread is enough
for all connections!

Reference: http://berb.github.io/diploma-thesis/original/042_serverarch.html

What is Node.js: Non-blocking I/O

- Also called **Asynchronous I/O**
- You are familiar with **blocking I/O** already...

For those who have taken CSCI/CENG 3150...

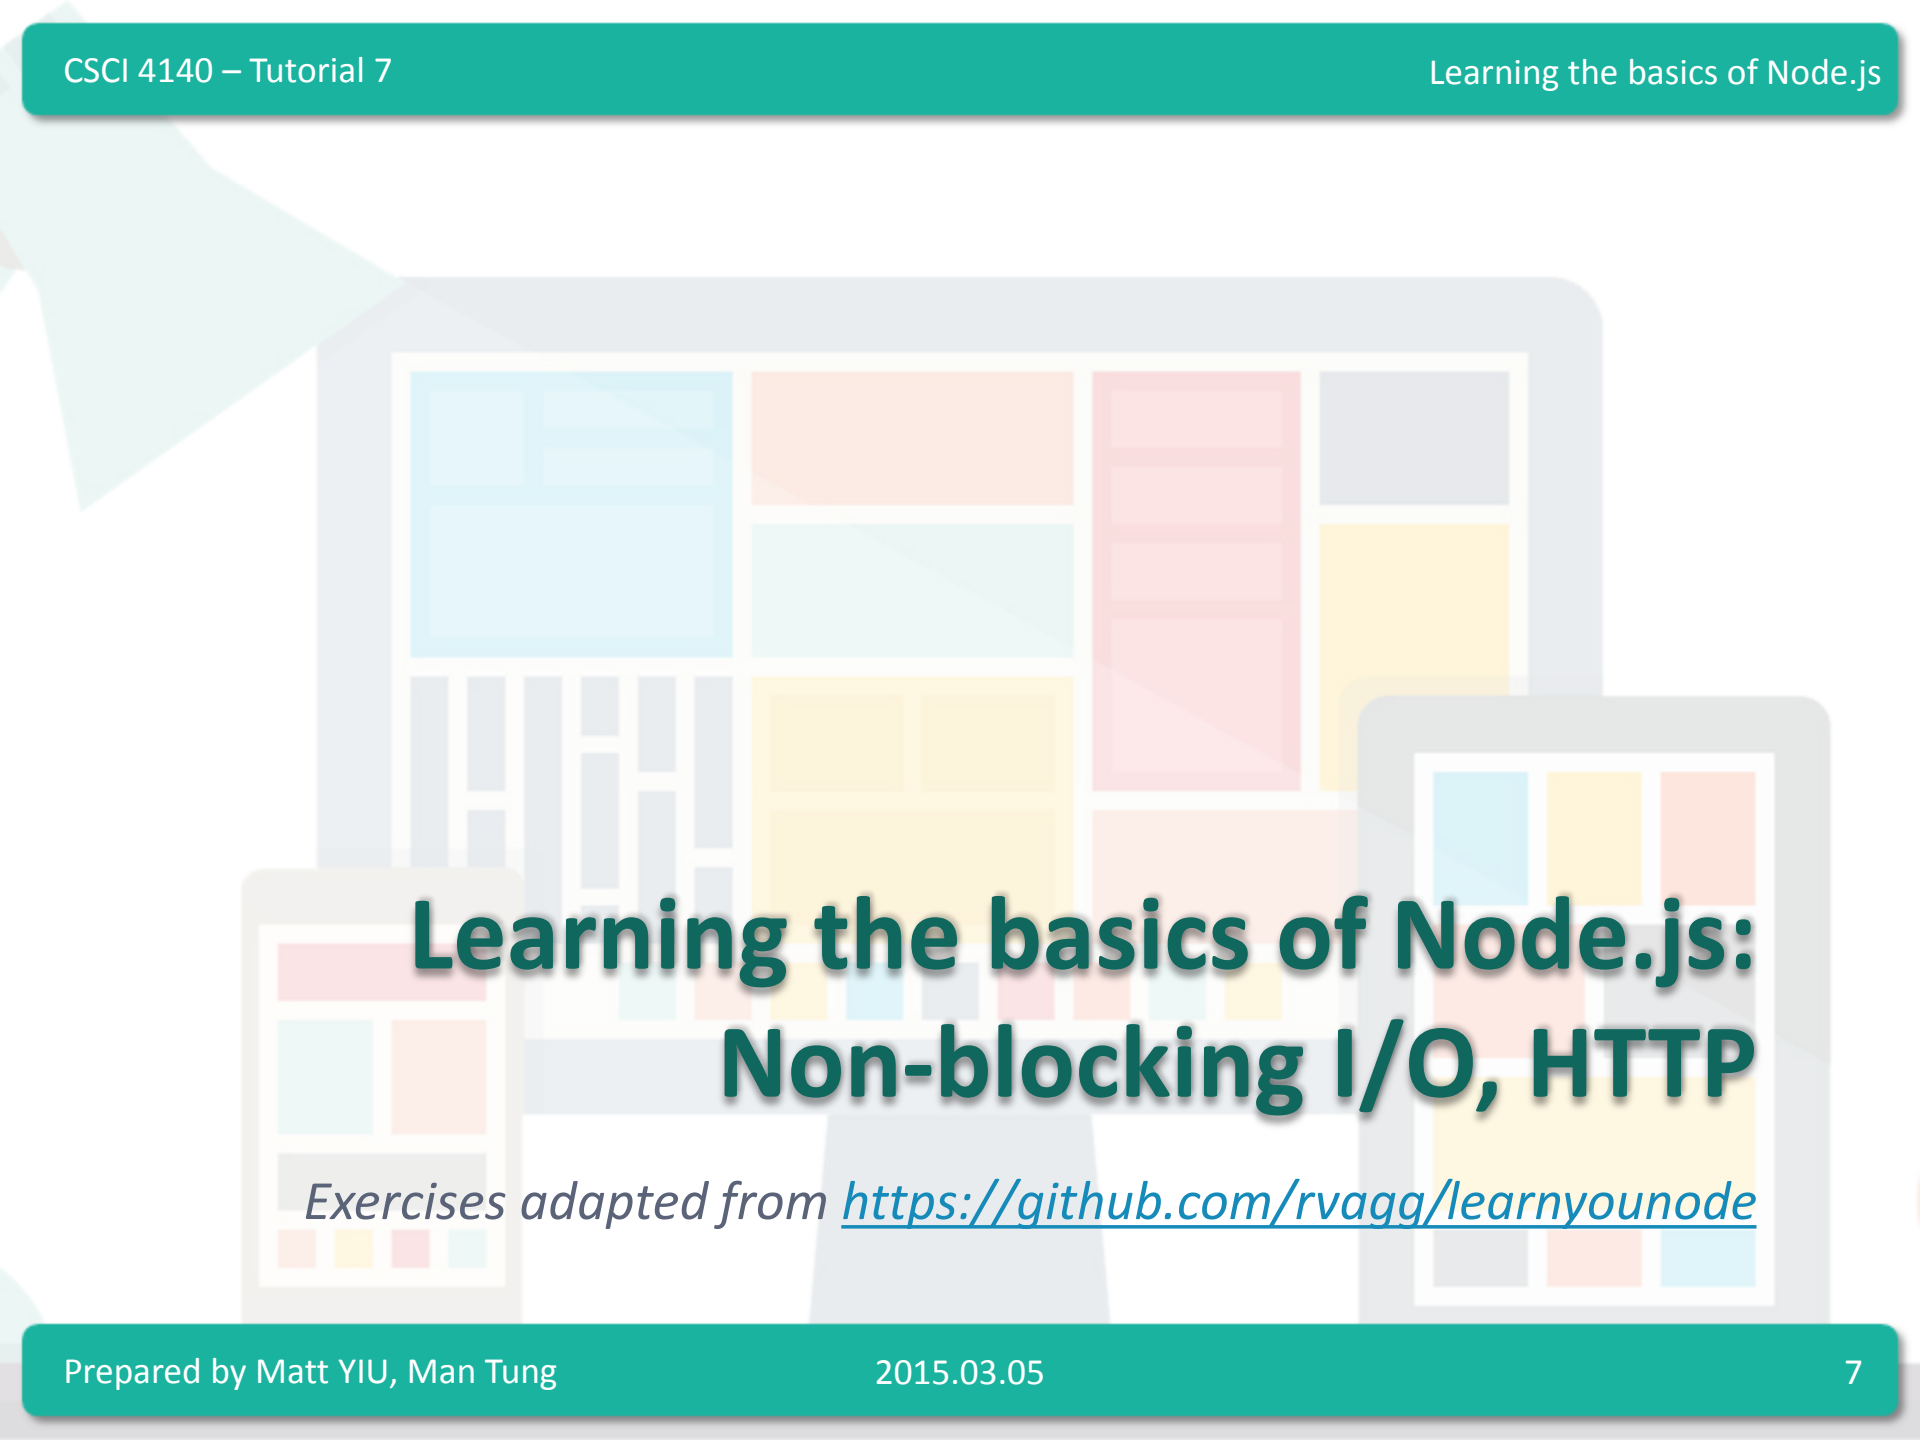


Node.js HTTP server

- HTTP is a **first class citizen** in Node
 - Forget about Apache / IIS / Nginx
- Say “Hello World!” with Node.js HTTP server:
 - Execute “**node nodejs/server.js**” in your terminal and visit <http://127.0.0.1:4140/> in your browser

```
var http = require( 'http' );  
http.createServer( function( request, response ) {  
    response.writeHead( 200, { 'Content-Type' : 'text/plain' } );  
    response.end( 'Hello World!\n' );  
} ).listen( 4140, '127.0.0.1' );  
  
console.log( 'Server running at http://127.0.0.1:4140/' );
```

nodejs/server.js

The background features a stylized illustration of a laptop and a tablet. The laptop screen is filled with various colored rectangular blocks in shades of blue, orange, red, and yellow, arranged in a grid-like pattern. The tablet, positioned in front of the laptop, also displays a similar arrangement of colored blocks. The overall aesthetic is clean and modern, with a focus on geometric shapes and a vibrant color palette.

Learning the basics of Node.js: Non-blocking I/O, HTTP

Exercises adapted from <https://github.com/rvagg/learnyounode>

Exercise 1: Hello World

- *Let's learn Node.js by doing exercises!*
- **Problem:** Write a program that prints the text “HELLO WORLD” to the console (stdout)
- Use the **console** API: <http://nodejs.org/api/console.html>

```
console.log( "HELLO WORLD" );
```

```
nodejs/ex1-hello.js
```

```
$ node nodejs/ex1-hello.js
```

Terminal

- Useful for debugging
 - Obviously you cannot call “alert()”...

Exercise 2: Baby steps

- **Problem:** Write a program that accepts one or more numbers as command-line arguments and prints the sum of those numbers to the console (stdout)
- Access **command-line arguments** from the **argv** property of the global **process** object
 - For example, executing “**node program.js 1 2 3**”

```
program.js console.log(process.argv);
```

Output ['node', '/home/mtyiou/program.js', '1', '2', '3']

- Note that the command-line arguments are strings
 - Convert the string into number with “**Number(<string>)**”

Exercise 3: My first I/O

- **Problem:** Write a program that uses a single **synchronous** filesystem operation to read a file and print the **number of newlines** it contains to the console (stdout), similar to running **cat file | wc -l**.
- We need the **fs** module from the Node core library
 - <http://nodejs.org/api/fs.html>
 - Load the fs module into a **variable**: `var fs = require('fs');`
- All synchronous (or blocking) filesystem methods end with “Sync”, e.g., “**fs.readFileSync(<file path>)**”
 - This method returns a **Buffer** object containing the complete contents of the file

Exercise 3: My first I/O

- **Buffer** objects are Node's way of efficiently representing arbitrary arrays of data
 - To convert them to strings, call “**toString()**” method on them, e.g.,
var str = buf.toString()
- To count the number of newlines in a string, you can split it using the “**split()**” method with the “**\n**” character as the delimiter
- Remember that the last line of the input file does not contain a newline

Exercise 4: My first asynchronous I/O

- **Problem:** Write a program that uses a single **asynchronous** filesystem operation to read a file and print the number of **newlines** it contains to the console (**stdout**), similar to running **cat file | wc -l**.
- **fs.readFile()** is the asynchronous version of **fs.readFileSync()**
 - This method returns **without blocking**
 - To read the file contents, you need to pass a **callback function** which will be called when the I/O completes
 - This concept is **extremely important** in **JavaScript** programming!

Exercise 4: My first asynchronous I/O

- The callback functions should have the following signatures:

```
function callback ( err, data ) { /* ... */ }
```

Represent an error

The Buffer object / string
containing the file contents

OR

```
function callback ( err, options, data ) { /* ... */ }
```

Pass "utf8" for the options argument to
get a string instead of an Buffer object

Exercise 5: Filtered 1s

- **Problem:** Create a program that prints a list of files in a given directory to the console using **asynchronous I/O**, filtered by the **extension** of the files
 - **1st argument:** A directory name
 - **2nd argument:** A file extension to filter by
- Similar to Exercise 4, but with **fs.readdir()**
 - http://nodejs.org/api/fs.html#fs_fs_readdir_path_callback
- You will also need **path.extname()** in the **path** module
 - http://nodejs.org/api/path.html#path_path_extname_p

Exercise 6: Make it modular

- **Problem:** Same as Exercise 5, but you need to make it **modular**
- Write a **module file** to do most of the work
 - The module should **export a function** which takes 3 arguments:
 1. The directory name
 2. The filename extension string (identical to the corresponding command-line argument)
 3. A callback function
 - The callback function should use the idiomatic node(**err, data**) convention
 - **err** is null if there is no errors; return the errors from `fs.readdir()` otherwise
 - **data** is the filtered list of files, as an Array
 - Nothing should be printed from your module file
 - Only print from the original program

Exercise 6: Make it modular

- From the problem statement, we induce the four requirements of a module:
 - **Export** a single function that takes exactly the **arguments** described
 - Call the **callback** exactly **once** with an error or some data as described
 - **Don't change anything** else, like global variables or stdout
 - Handle all the **errors** that may occur and pass them to the callback
 - Do **early-returns** within callback functions if there is an error
- A good Node.js developer should follow these rules!

Exercise 6: Make it modular

- In the module file (e.g., `module.js`), assign a function to the **`module.exports`** object to define a **single function export**:

```
module.exports = function (args) { /* ... */ }
```

- In your program, load the module (`module.js`) using the **`require()`** call ("`./`" indicates that it is a local module):

```
var module = require( './module' );
```

- **Note:** "`.js`" can be omitted
- The `require()` call returns what you export in the module file
 - In this example, it returns a function that you can call directly!

Exercise 7: HTTP client

- **Problem:** Write a program that performs an **HTTP GET request** to a URL provided to you as the first command-line argument. Write the String contents of each “data” event from the response to a new line on the console (stdout).
 - **Note:** There is a sample scenario in Assignment 2 – retrieving video title from YouTube server using an HTTP GET request
- Use the **http.get()** method in the **http** module
 - http://nodejs.org/api/http.html#http_http_get_options_callback
 - **1st argument:** The URL you want to GET
 - **2nd argument:** A callback with the following signature:

```
function callback ( response ) { /* ... */ }
```

Exercise 7: HTTP client

- The **response** object is a **Node Stream** object
 - It is an object that emits events
 - Register an **event listener** (**.on(*, callback)**) to handle the event
 - This is the core of “**event-driven architecture**”
 - For `http.get()`, the three events that are of most interests are: “**data**”, “**error**” and “**end**”
 - See http://nodejs.org/api/http.html#http_http_incomingmessage and http://nodejs.org/api/stream.html#stream_class_stream_readable
- The response object has a **setEncoding()** method
 - If you call this method with “utf8”, the **data** events emit Strings instead of the standard Node **Buffer** objects

Exercise 8: HTTP collect

- **Problem:** Write a program that performs an HTTP GET request to a URL provided to you as the first command-line argument. Collect **all data** from the server (not just the first “data” event) and then write two lines to the console (stdout).
 - **1st line:** The **number of characters** received from the server
 - **2nd line:** The **complete String of characters** sent by the server
- Two approaches:
 - Collect and append data across multiple “data” events. Write the output when an “end” event is emitted
 - Use a **third-party package** to abstract the difficulties involved in collecting an entire stream of data, e.g., **bl** and **concat-stream**

Exercise 8: HTTP collect

- Let's try the second approach to explore an important component in Node.js: **npm** – the **package manager** for node
 - FYI, the package manager for Python is **pip**
- To install the Node package **bl**, type in the terminal:

```
$ npm install bl
```

 - npm will download and install the latest version of the package into a subdirectory named **node_modules**
- When you write “**var bl = require('bl');**” in your program, Node will first look in the **core modules**, and then in the **node_modules** directory where the package is located.
- Read <https://www.npmjs.com/package/bl> for its usage

Exercise 9: Juggling async

- **Problem:** Same as Exercise 8, but this time you will be provided with **3 URLs** as the first 3 command-line arguments
 - Print the complete content provided by each of the URLs to the console (stdout), one line per URL
 - No need to print out the length
 - The content must be printed out in the **same order** as the URLs are provided to you as command-line arguments
- **This exercise is tricky!**
 - **`http.get()`** is an asynchronous call
 - The callback function is executed when any of the servers response
 - The responses will probably be **out of order!**
 - You need to **queue the results** and print the data when all data is ready

Exercise 10: Time server

- **Problem:** Write a TCP time server!
 - Your server should listen to **TCP connections** on the port provided by the first argument to your program
 - For each connection you must write the **current date & 24 hour time** in the format: “**YYYY-MM-DD hh:mm**”, followed by a newline character
 - Month, day, hour and minute must be zero-filled to 2 integers
 - For example: “2013-07-06 17:42”
- This exercise demonstrates the power of Node.js!
 - Challenge to CSCI 4430 students: Solve this problem in **C/C++ socket programming**!

Exercise 10: Time server

- To create a **raw TCP server**, use the **net** module
 - Use the method named **net.createServer()**
 - It returns an instance of your server
 - To start listening on a particular port, use **server.listen(<port>)**
 - It takes a callback function with the following signature:

```
function callback ( socket ) { /* ... */ }
```

- The **socket** object passed into the callback function contains a lot of metadata regarding the connection
- To write data to the **socket**: `socket.write(data);`
- To close the **socket**: `socket.end();`
- Ref.: <http://nodejs.org/api/net.html>

Can be
combined

`socket.end(data);`

Exercise 10: Time server

- To create the date, you will need to create a custom format from a **new Date()** object
- The following methods will be useful:
 - `date.getFullYear()`
 - `date.getMonth()` // starts at 0
 - `date.getDate()` // returns the day of month
 - `date.getHours()`
 - `date.getMinutes()`

Exercise 11: HTTP file server

- *Now we are ready to learn how to use Node.js to implement server-side program!*
- **Problem:** Write an **HTTP server** that **serves the same text file** for each request it receives
 - **1st argument:** **Port number** that the server listens on
 - **2nd argument:** The **location** of the file to serve
- You must use the **fs.createReadStream()** method to stream the file contents to the response
 - It creates a **stream** representing the file
 - Use **src.pipe(dst)** to pipe data from the **src** stream to the **dst** stream

Exercise 11: HTTP file server

- Use the **http** module to create an HTTP server
 - **http.createServer()** take a callback that is called once for each connection received by your server

```
function callback ( request, response ) { /* ... */ }
```

- The two arguments are **Node stream** objects representing the HTTP request and the corresponding response
 - Request is used for **fetch properties**, e.g., the header and query string
 - Response is for **sending data to the client**, both headers and body
- **Ref.:** <http://nodejs.org/api/http.html>

Exercise 12: HTTP upercaserer

- **Problem:** Write an HTTP server that **receives only POST requests** and converts incoming **POST body characters** to **upper-case** and returns it to the client
 - **1st argument:** **Port number** that the server listens on
- You can use the “**through2-map**” module to create a **transform stream** using only a single function that takes a **chunk of data** and returns a chunk of data
 - Install **through2-map** using **npm**
 - Read <https://www.npmjs.com/package/through2-map> for its usage

Exercise 13: HTTP JSON API server

- **Problem:** Write an HTTP server that serves **JSON data** when it
 - Receives a GET request to the path **“/api/parsetime”**
 - The JSON response should contain only **‘hour’**, **‘minute’** and **‘second’** properties
 - Receives a GET request to the path **“/api/unixtime”**
 - The JSON response should contain the **UNIX epoch time in milliseconds** (the number of milliseconds since 1 Jan 1970 00:00:00 UTC) under the property **‘unixtime’**
 - Both requests accept a query string with a key **‘iso’** and an ISO-format time as the value
 - **1st argument of the program:** **Port number** that the server listens on

Exercise 13: HTTP JSON API server

- Use the **`url.parse()`** method in the **`url`** module to parse the URL and query string
 - *Ref.:* <http://nodejs.org/api/url.html>
- Use **`JSON.stringify()`** to convert an object into JSON string format
- To parse a date in **ISO format**, use **`new Date(<ISO date string>)`**
- Use **`date.getTime()`** to get the **UNIX epoch time in milliseconds**

Congratulations!

- You have learnt the fundamental concepts involved in Node.js development!
- To develop web applications even faster, we will use a **web framework** called **Express**
 - Please refer to the corresponding tutorial slides

– End –